

End-to-End ETL Pipeline for Data Analysis and Visualization

Authors: Aniket Gulab Khedkar

Contact: aniketgk@umich.edu

0) Stack

- **Source:** Kaggle Olist CSVs
- **DB:** Microsoft SQL Server (on-prem or Azure SQL)
- **Ingest:** Python (pandas + SQLAlchemy + pyodbc, fast_executemany=True)
- **Modeling:** dbt-sqlserver (T-SQL) to build dimensions & facts
- **Quality:** lightweight checks (Python)
- **Viz:** Power BI / Tableau (+ optional Streamlit app)

1) Project Layout

```
olist-mssql-etl/
├── configs/
│   ├── .env.example
│   └── dbt_profile.yml
├── sql/
│   ├── create_schemas.sql
│   ├── create_staging_tables.sql
│   ├── helper_indexes.sql
│   └── dbt_project/
│       ├── dbt_project.yml
│       └── models/
│           ├── sources.yml
│           └── staging/
```

```
├── stg_orders.sql
├── stg_order_items.sql
├── stg_payments.sql
├── stg_customers.sql
├── stg_products.sql
├── stg_sellers.sql
├── stg_geolocation.sql
├── stg_reviews.sql
├── warehouse/
│   ├── dim_date.sql
│   ├── dim_customer.sql
│   ├── dim_product.sql
│   └── dim_seller.sql
├── fact_order_items.sql
├── fact_payments.sql
├── fact_orders.sql
├── marts/
├── sales_summary.sql
├── scripts/
│   ├── kaggle_download.sh
│   ├── load_csvs_to_staging.py
│   ├── run_quality_checks.py
│   └── streamlit_app.py
├── requirements.txt
└── README.md
```

2) Install & Configure

2.1 Python env

```
python -m venv .venv && source
.venv/bin/activate # (Windows:
.venv\Scripts\activate)
pip install -r requirements.txt
```

requirements.txt

```
python-dotenv>=1.0
pandas>=2.2
pyodbc>=5.1
SQLAlchemy>=2.0
tqdm>=4.66
# dbt for SQL Server
dbt-core>=1.7
dbt-sqlserver>=1.7
# dashboard
streamlit>=1.37
plotly>=5.22
```

Install ODBC Driver 18 for SQL Server on your OS.

2.2 Kaggle CLI

- Create/deploy your Kaggle API token: <https://www.kaggle.com/settings/account>
- Place kaggle.json under ~/.kaggle/ (Linux/Mac) or %USERPROFILE%\kaggle\ (Windows).

2.3 .env (copy from configs/.env.example)

```
MSSQL_SERVER=localhost
MSSQL_PORT=1433
MSSQL_DATABASE=olistdb
MSSQL_USERNAME=sa
MSSQL_PASSWORD=YourStrong!Passw0rd
MSSQL_DRIVER=ODBC Driver 18 for
SQL Server
```

3) Download Olist CSVs (Kaggle)

scripts/kaggle_download.sh

```
#!/usr/bin/env bash
set -euo pipefail
mkdir -p data/raw
kaggle datasets download -d
olistbr/brazilian-ecommerce -p
data/raw
unzip -o data/raw/brazilian-
ecommerce.zip -d data/raw
```

Run:

```
bash scripts/kaggle_download.sh
```

4) Create Schemas & Staging Tables (SQL Server)

4.1 Schemas

```
-- sql/create_schemas.sql
USE OlistDB;
```

```
IF NOT EXISTS (SELECT * FROM
sys.schemas WHERE name='stg')
EXEC('CREATE SCHEMA stg');
IF NOT EXISTS (SELECT * FROM
sys.schemas WHERE name='wh')
EXEC('CREATE SCHEMA wh');
```

4.2 Staging tables (1:1 with Kaggle CSVs)

```
-- sql/create_staging_tables.sql
USE OlistDB;
```

```
CREATE TABLE stg.olist_orders (
    order_id
    VARCHAR(50) PRIMARY KEY,
    customer_id
    VARCHAR(50),
    order_status
    VARCHAR(32),
    order_purchase_timestamp
    DATETIME2,
    order_approved_at
    DATETIME2,
    order_delivered_carrier_date
    DATETIME2,
    order_delivered_customer_date
    DATETIME2,
    order_estimated_delivery_date
    DATETIME2
);
```

```
CREATE TABLE stg.olist_order_items
(
    order_id          VARCHAR(50),
    order_item_id     INT,
    product_id        VARCHAR(50),
    seller_id         VARCHAR(50),
    shipping_limit_date DATETIME2,
    price             DECIMAL(18,2),
    freight_value     DECIMAL(18,2)
);
```

```
CREATE TABLE
stg.olist_order_payments (
    order_id          VARCHAR(50),
    payment_sequential INT,
    payment_type      VARCHAR(32),
    payment_installments INT,
    payment_value     DECIMAL(18,2)
);
```

```
CREATE TABLE stg.olist_customers (
    customer_id          VARCHAR(50)
    PRIMARY KEY,
    customer_unique_id   VARCHAR(50),
    customer_zip_code_prefix INT,
    customer_city
    NVARCHAR(100),
    customer_state       NVARCHAR(10)
```

```

);

CREATE TABLE stg.olist_products (
    product_id
    VARCHAR(50) PRIMARY KEY,
    product_category_name
    NVARCHAR(100),
    product_name_length
    INT,
    product_description_length
    INT,
    product_photos_qty
    INT,
    product_weight_g
    INT,
    product_length_cm
    INT,
    product_height_cm
    INT,
    product_width_cm
    INT
);

CREATE TABLE stg.olist_sellers (
    seller_id
    VARCHAR(50) PRIMARY KEY,
    seller_zip_code_prefix INT,
    seller_city
    NVARCHAR(100),
    seller_state
    NVARCHAR(10)
);

CREATE TABLE stg.olist_geolocation
(
    geolocation_zip_code_prefix INT,
    geolocation_lat FLOAT,
    geolocation_lng FLOAT,
    geolocation_city NVARCHAR(100),
    geolocation_state NVARCHAR(10)
);

CREATE TABLE
stg.product_category_name_translat
ion (
    product_category_name
    NVARCHAR(100),
    product_category_name_english
    NVARCHAR(100)
);

CREATE TABLE
stg.olist_order_reviews (
    review_id
    VARCHAR(50)
    PRIMARY KEY,
    order_id
    VARCHAR(50),
    review_score
    INT,
    review_comment_title
    NVARCHAR(255),
    review_comment_message
    NVARCHAR(MAX),
    review_creation_date DATETIME2,

```

```

    review_answer_timestamp
    DATETIME2
);

```

4.3 Helpful indexes

```

-- sql/helper_indexes.sql
CREATE INDEX IX_orders_purchase_ts
ON
stg.olist_orders(order_purchase_ti
mestamp);
CREATE INDEX IX_items_order ON
stg.olist_order_items(order_id);
CREATE INDEX IX_payments_order ON
stg.olist_order_payments(order_id)
;
CREATE INDEX IX_customers_unique
ON
stg.olist_customers(customer_uniqu
e_id);
CREATE INDEX IX_geo_zip ON
stg.olist_geolocation(geolocation_
zip_code_prefix);

```

5) Load CSVs → Staging (Python bulk upsert)

scripts/load_csvs_to_staging.py

```

import os, glob, pandas as pd
from sqlalchemy import
create_engine, text
from dotenv import load_dotenv

load_dotenv()
SERVER=os.getenv('MSSQL_SERVER');
PORT=os.getenv('MSSQL_PORT','1433'
)
DB=os.getenv('MSSQL_DATABASE');
USER=os.getenv('MSSQL_USERNAME');
PWD=os.getenv('MSSQL_PASSWORD')
DRV=os.getenv('MSSQL_DRIVER','ODBC
Driver 18 for SQL Server')

odbc=f"DRIVER={{DRV}};SERVER={SE
RVER},{PORT};DATABASE={DB};UID={US
ER};PWD={PWD};Encrypt=yes;TrustSer
verCertificate=yes;".replace('
','+')
engine=create_engine(f"mssql+pyodb
c:///odbc_connect={odbc}",
fast_executemany=True)

CSV_MAP={
    'olist_orders_dataset.csv':
    ('stg.olist_orders', {

```

```
'order_purchase_timestamp':'datetime64[ns]',
```

```
'order_approved_at':'datetime64[ns]',
'order_delivered_carrier_date':'datetime64[ns]',
```

```
'order_delivered_customer_date':'datetime64[ns]',
'order_estimated_delivery_date':'datetime64[ns]')),
'olist_order_items_dataset.csv':
('stg.olist_order_items',
{'shipping_limit_date':'datetime64[ns]'}),
```

```
'olist_order_payments_dataset.csv':
('stg.olist_order_payments', {}),
'olist_customers_dataset.csv':
('stg.olist_customers', {}),
'olist_products_dataset.csv':
('stg.olist_products', {}),
'olist_sellers_dataset.csv':
('stg.olist_sellers', {}),
'olist_geolocation_dataset.csv':
('stg.olist_geolocation', {}),
```

```
'olist_order_reviews_dataset.csv':
('stg.olist_order_reviews', {
```

```
'review_creation_date':'datetime64[ns]',
'review_answer_timestamp':'datetime64[ns]'}),
```

```
'product_category_name_translation.csv':
('stg.product_category_name_translation', {})
}
```

```
with engine.begin() as con:
    for fname, (table, dtypes) in CSV_MAP.items():
```

```
path=os.path.join('data', 'raw', fname)
```

```
df=pd.read_csv(path,
dtype=str, keep_default_na=False)
# Cast specific dtypes
for col,dt in dtypes.items():
    if col in df.columns:
```

```
df[col]=pd.to_datetime(df[col],
errors='coerce')
# Numeric casts
for c in ['order_item_id', 'price', 'freight_value',
'payment_sequential', 'payment_installments', 'payment_value',
```

```
'product_name_length', 'product_description_length',
'product_photos_qty',
```

```
'product_weight_g', 'product_length_cm',
'product_height_cm', 'product_width_cm',
```

```
'customer_zip_code_prefix', 'seller_zip_code_prefix',
'review_score',
```

```
'geolocation_zip_code_prefix', 'geolocation_lat',
'geolocation_lng']:
    if c in df.columns:
```

```
df[c]=pd.to_numeric(df[c],
errors='coerce')
# Load: replace table on first run for idempotency
```

```
df.to_sql(table.split('.')[-1], con,
schema=table.split('.')[0], if_exists='append',
index=False)
print(f"Loaded {len(df):,} rows into {table}")
```

Run:

```
python
scripts/load_csvs_to_staging.py
```

6) Transform → Warehouse (dbt-sqlserver)

6.1 Configure dbt profile

configs/dbt_profile.yml

```
default:
  target: dev
  outputs:
    dev:
      type: sqlserver
      driver: 'ODBC Driver 18 for SQL Server'
      server: '{{
env_var("MSSQL_SERVER") }}'
      port: 1433
      user: '{{
env_var("MSSQL_USERNAME") }}'
      password: '{{
env_var("MSSQL_PASSWORD") }}'
      database: '{{
env_var("MSSQL_DATABASE") }}'
      schema: 'wh'
```

```

encrypt: true
trust_cert: true

```

6.2 Declare sources

dbt_project/models/sources.yml

```

version: 2
sources:
  - name: stg
    schema: stg
    tables:
      - name: olist_orders
      - name: olist_order_items
      - name: olist_order_payments
      - name: olist_customers
      - name: olist_products
      - name: olist_sellers
      - name: olist_geolocation
      - name: olist_order_reviews
      - name:
product_category_name_translation

```

6.3 Staging views (standardize types)

dbt_project/models/staging/stg_orders.sql

```

select
  order_id,
  customer_id,
  order_status,
  cast(order_purchase_timestamp as
datetime2) as order_purchase_ts,

cast(order_estimated_delivery_date
as date) as estimated_delivery_dt,

cast(order_delivered_customer_date
as datetime2) as
delivered_customer_ts
from {{
source('stg','olist_orders') }}

```

dbt_project/models/staging/stg_order_items.sql

```

select
  order_id,
  order_item_id,
  product_id,
  seller_id,
  cast(shipping_limit_date as
datetime2) as shipping_limit_ts,
  cast(price as decimal(18,2)) as
price,

```

```

  cast(freight_value as
decimal(18,2)) as freight_value
from {{
source('stg','olist_order_items')
}}

```

dbt_project/models/staging/stg_payments.sql

```

select
  order_id,
  payment_type,
  payment_installments,
  cast(payment_value as
decimal(18,2)) as payment_value
from {{
source('stg','olist_order_payments')
}}

```

dbt_project/models/staging/stg_customers.sql

```

select
  customer_id,
  customer_unique_id,
  customer_zip_code_prefix,
  customer_city,
  customer_state
from {{
source('stg','olist_customers') }}

```

dbt_project/models/staging/stg_products.sql

```

select p.product_id,

coalesce(t.product_category_name_english, p.product_category_name)
as product_category,
  p.product_weight_g,
p.product_length_cm,
p.product_height_cm,
p.product_width_cm
from {{
source('stg','olist_products') }}
p
left join {{
source('stg','product_category_name_translation') }} t
  on t.product_category_name =
p.product_category_name

```

dbt_project/models/staging/stg_sellers.sql

```
select seller_id,
seller_zip_code_prefix,
seller_city, seller_state
from {{
source('stg','olist_sellers') }}
```

dbt_project/models/staging/stg_reviews.sql

```
select order_id, review_score
from {{
source('stg','olist_order_reviews'
) }}
```

6.4 Warehouse models (star schema)

dbt_project/models/warehouse/dim_date.sql

```
with d as (
    select cast(dateadd(day,
v.number, '2016-01-01') as date)
as d
    from master..spt_values v where
v.type='P' and v.number between 0
and 1826 -- 5 years
)
select cast(format(d,'yyyyMMdd')
as int) as date_key,
    d as date_value,
    year(d) as year,
    datepart(quarter,d) as
quarter,
    month(d) as month,
    day(d) as day,
    datepart(week,d) as
week_of_year
```

dbt_project/models/warehouse/dim_customer.sql

```
select
    c.customer_unique_id as
customer_key,
    any_value(c.customer_id) as
sample_customer_id,
    c.customer_city,
    c.customer_state
from {{ ref('stg_customers') }} c
group by c.customer_unique_id,
c.customer_city, c.customer_state
```

dbt_project/models/warehouse/dim_product.sql

```
select
    p.product_id as product_key,
    p.product_category,
    p.product_weight_g,
    p.product_length_cm,
    p.product_height_cm,
    p.product_width_cm
from {{ ref('stg_products') }} p
```

dbt_project/models/warehouse/dim_seller.sql

```
select
    s.seller_id as seller_key,
    s.seller_city,
    s.seller_state
from {{ ref('stg_sellers') }} s
```

dbt_project/models/warehouse/fact_orders.sql

```
select
    o.order_id,
    dc.customer_key,

cast(format(o.order_purchase_ts,'y
yyyMMdd') as int) as date_key,
    o.order_status,
    o.estimated_delivery_dt,
    o.delivered_customer_ts,
    i.total_items,
    i.items_value,
    i.freight_value,
    p.payments_value,
    r.avg_review_score
from {{ ref('stg_orders') }} o
left join (
    select order_id,
        sum(price) as
items_value,
        sum(freight_value) as
freight_value,
        count(*) as total_items
    from {{ ref('stg_order_items')
}}
    group by order_id
) i on i.order_id = o.order_id
left join (
    select order_id,
sum(payment_value) as
payments_value
    from {{ ref('stg_payments') }}
    group by order_id
) p on p.order_id = o.order_id
left join (
    select order_id,
avg(review_score*1.0) as
avg_review_score
    from {{ ref('stg_reviews') }}
```

```

    group by order_id
) r on r.order_id = o.order_id
left join (
    select customer_id,
    customer_unique_id from {{
    ref('stg_customers') }}
) cu on cu.customer_id =
o.customer_id
left join {{ ref('dim_customer')
}} dc on dc.customer_key =
cu.customer_unique_id

```

dbt_project/models/warehouse/fact_order_items.sql

```

select
    oi.order_id,

cast(format(o.order_purchase_ts,'y
yyyMMdd') as int) as date_key,
    dp.product_key,
    ds.seller_key,
    oi.order_item_id,
    oi.price,
    oi.freight_value
from {{ ref('stg_order_items') }}
oi
left join {{ ref('stg_orders') }}
o on o.order_id = oi.order_id
left join {{ ref('dim_product') }}
dp on dp.product_key =
oi.product_id
left join {{ ref('dim_seller') }}
ds on ds.seller_key = oi.seller_id

```

dbt_project/models/marts/sales_summary.sql

```

select
    dd.year,
    dd.month,
    dp.product_category,
    sum(fo.items_value) as revenue,
    sum(fo.freight_value) as
freight,
    sum(fo.total_items) as items,
    avg(fo.avg_review_score) as
avg_review
from {{ ref('fact_orders') }} fo
join {{ ref('dim_date') }} dd on
dd.date_key = fo.date_key
join {{ ref('dim_product') }} dp
on dp.product_key in (
    select product_key from {{
    ref('fact_order_items') }} foi
where foi.order_id = fo.order_id
)
group by dd.year, dd.month,
dp.product_category

```

Run dbt:

```

export DBT_PROFILES_DIR=./configs
cd dbt_project && dbt run && dbt
test

```

7) Quick Data Quality Checks (Python)

scripts/run_quality_checks.py

```

from sqlalchemy import
create_engine
import pandas as pd, os
from dotenv import load_dotenv
load_dotenv()
engine=create_engine(

f"mssql+pyodbc:///odbc_connect=DR
IVER={{ODBC Driver 18 for SQL
Server}};SERVER={{os.getenv('MSSQL_
SERVER')}};{os.getenv('MSSQL_PORT',
'1433')}};DATABASE={{os.getenv('MSSQ
L_DATABASE')}};UID={{os.getenv('MSSQ
L_USERNAME')}};PWD={{os.getenv('MSSQ
L_PASSWORD')}};Encrypt=yes;TrustSer
verCertificate=yes;".replace('
','+')
)

checks={

'stg.olist_orders':['order_id','cu
stomer_id','order_purchase_timesta
mp'],

'stg.olist_order_items':['order_id
','order_item_id','product_id','pr
ice'],

'stg.olist_order_payments':['order
_id','payment_value'],
}

with engine.begin() as con:
    for table, cols in
checks.items():
        df=pd.read_sql(f"select * from
{table}", con)
        missing=[c for c in cols if c
not in df.columns or
df[c].isna().any()]
        if missing: raise
AssertionError(f"{table} failed
not-null check on {missing}")
print("Quality checks passed")

```

Run:

```
python
scripts/run_quality_checks.py
```

8) Dashboards

Power BI (recommended)

1. Connect → SQL Server → Database: **OlistDB**.
2. Import: wh.dim_date, wh.dim_customer, wh.dim_product, wh.dim_seller, wh.fact_orders, wh.fact_order_items, and view wh.sales_summary (if materialized as view).
3. Relationships:
 - o fact_orders.date_key → dim_date.date_key
 - o fact_order_items.date_key → dim_date.date_key
 - o fact_order_items.product_key → dim_product.product_key
 - o fact_order_items.seller_key → dim_seller.seller_key

DAX Measures

```
Total Revenue :=
SUMX(wh_fact_orders,
wh_fact_orders[items_value])
Total Freight :=
SUM(wh_fact_orders[freight_value])
Total Orders :=
DISTINCTCOUNT(wh_fact_orders[order_id])
Total Items :=
SUM(wh_fact_orders[total_items])
AOV := DIVIDE([Total Revenue],
[Total Orders])
OnTime Delivery % := DIVIDE(
    CALCULATE([Total Orders],
    FILTER(wh_fact_orders,
wh_fact_orders[delivered_customer_ts] <=
wh_fact_orders[estimated_delivery_dt])),
[Total Orders])
)
```

```
Avg Review :=
AVERAGE(wh_fact_orders[avg_review_score])
Repeat Customer % :=
VAR t = SUMMARIZE(wh_dim_customer,
wh_dim_customer[customer_key],
"orders",
CALCULATE(DISTINCTCOUNT(wh_fact_orders[order_id])))
RETURN DIVIDE(COUNTROWS(FILTER(t,
[orders] > 1)), COUNTROWS(t))
```

Streamlit (optional)

scripts/streamlit_app.py

```
import os, pandas as pd,
plotly.express as px
from sqlalchemy import
create_engine
import streamlit as st

st.set_page_config(page_title="Olist Sales Analytics",
layout="wide")
engine = create_engine(

f"mssql+pyodbc:///odbc_connect=DRIVER={{ODBC Driver 18 for SQL Server}};SERVER={os.getenv('MSSQL_SERVER')},{os.getenv('MSSQL_PORT','1433')};DATABASE={os.getenv('MSSQL_DATABASE')};UID={os.getenv('MSSQL_USERNAME')};PWD={os.getenv('MSSQL_PASSWORD')};Encrypt=yes;TrustServerCertificate=yes;".replace(' ','+')
)

@st.cache_data
def load(sql):
    with engine.begin() as con:
        return pd.read_sql(sql, con)

df = load("select * from wh.sales_summary")
st.title("Olist - Sales Summary")
col1, col2 = st.columns(2)
col1.plotly_chart(px.line(df, x="month", y="revenue", color="year", title="Revenue by Month"), use_container_width=True)
col2.plotly_chart(px.bar(df, x="product_category", y="revenue", title="Revenue by Category"), use_container_width=True)
```

Run:


```
streamlit run
scripts/streamlit_app.py
```

- Adjust indexes for large deployments; partition by date if needed.
-

9) End-to-End Runbook

```
# 1) env & deps
python -m venv .venv && source
.venv/bin/activate
pip install -r requirements.txt

# 2) download Kaggle data
bash scripts/kaggle_download.sh

# 3) create schemas + staging
tables in SQL Server
sqlcmd -S localhost -d OlistDB -U
sa -P YourStrong!Passw0rd -i
sql/create_schemas.sql
sqlcmd -S localhost -d OlistDB -U
sa -P YourStrong!Passw0rd -i
sql/create_staging_tables.sql
sqlcmd -S localhost -d OlistDB -U
sa -P YourStrong!Passw0rd -i
sql/helper_indexes.sql

# 4) load CSVs to staging
python
scripts/load_csvs_to_staging.py

# 5) build warehouse with dbt
export DBT_PROFILES_DIR=./configs
cd dbt_project && dbt run && dbt
test && cd ..

# 6) (optional) quality checks
python
scripts/run_quality_checks.py

# 7) (optional) run Streamlit
dashboard
streamlit run
scripts/streamlit_app.py
```

10) Notes

- **Revenue** uses sum of `order_items.price` (payments can be used for cross-checks).
- **On-time delivery** compares actual delivered timestamp vs. estimated delivery date.
- **Repeat customers** defined via `customer_unique_id` (stable across orders).